



# Reducing Software Risk

**Rob Bocchino**

**NASA Jet Propulsion Laboratory**

**October 22, 2024**

Copyright © 2024 California Institute of Technology.  
Government sponsorship acknowledged.

Reference herein to any specific commercial product, process, or service by trade name, trademark, manufacturer, or otherwise, does not constitute or imply its endorsement by the United States Government or the Jet Propulsion Laboratory, California Institute of Technology.



**Jet Propulsion Laboratory**  
California Institute of Technology



# Reducing Software Risk

- Flight code must perform well and be reliable
- C and C++ are widely used
  - High performance
  - Direct interface to hardware
- But they come with risks
  - Code can have undefined or unexpected behavior
  - Compiler can catch some errors
  - But many it cannot (e.g., integer overflow, dangling pointers)
- This section of the course will show how to reduce the risk



# Outline

- **Check numeric bounds**
- Use bounded loops
- Use assertions
- Avoid dangerous code
- Run static analysis
- Follow a coding standard



# Check Numeric Bounds

- Integer and FP types in C and C++ are not numbers
- They are approximations of numbers with a finite representation
  - Integer types have a limited range
  - FP types have a limited range and a limited precision
- In FSW, you should avoid using platform-specific integer types
  - Instead of *int* and *unsigned*, use *int32\_t* and *uint32\_t*
  - This way the bounds are known and are the same across platforms
  - Use *int* if and only if it's required by a library or system interface
    - E.g., the return value of *main* is *int*
    - Many C library calls return *int* as their error codes



# Checking Integer Conversions

```
#include <cinttypes>

int32_t i64_to_i32(int64_t x) {
    return static_cast<int32_t>(x);
}
```

## Unchecked Conversion

```
#include <cinttypes>
#include <limits>

enum class Status { FAILURE, SUCCESS };

Status i64_to_i32(int64_t x, int32_t& result) {
    Status status = Status::FAILURE;
    if ((x >= std::limits<int32_t>::min()) &&
        (x <= std::limits<int32_t>::max())) {
        result = static_cast<int32_t>(x);
        status = Status::SUCCESS;
    }
    return status;
}
```

## Checked Conversion



# Checking Integer Overflow

```
int32_t add_i32(  
    int32_t a,  
    int32_t b  
) {  
    return a + b;  
}
```

## Unchecked Addition

```
Status add_i32(  
    int32_t a,  
    int32_t b,  
    int32_t& result  
) {  
    int64_t c = static_cast<int64_t>(a) +  
                static_cast<int64_t>(b);  
    return i64_to_i32(c, result);  
}
```

## Checked Addition

**Exercise:** Check the addition using only *int32\_t* variables  
**Hint:** Consider cases such as  $a \geq 0$  and  $b \geq 0$ ,  $a \leq 0$  and  $b \leq 0$ , etc.



# Beware of Floating Precision Loss

```
#include <stdio.h>

int
main(void)
{  float f, of, cnt, i;

    f = 1;

    for (cnt = 1; cnt < 47; cnt++)
    {
        f /= 10;

        for (of = f, i = 1; i <= cnt; i++)
        {
            of *= 10;
        }

        printf("%9.3g    %9.8f\n", f, of);
    }
}
```

0.1	1.00000000
0.01	0.99999994
0.001	0.99999994
0.0001	0.99999994
1e-05	0.99999994
1e-06	0.99999994
...	
1e-38	0.99999994
1e-39	1.00000024
1e-40	0.99999452
1e-41	0.99996656
1e-42	1.00052714
9.95e-44	0.99492157
9.81e-45	0.98090899
1.4e-45	1.40129852
0	0.00000000



# Outline

- Check numeric bounds
- **Use bounded loops**
- Use assertions
- Avoid dangerous code
- Run static analysis
- Follow a coding standard





# Use Bounded Loops

```
while (true) {  
    Queue::Item *item = nullptr;  
    const Status status = queue.pop(item);  
    if (status != Status::SUCCESS) {  
        break;  
    }  
    dispatch(item);  
}
```

**Unbounded Loop**

```
const size_t queueDepth = queue.getDepth();  
for (size_t i = 0; i < queueDepth; i++) {  
    Queue::Item *item = nullptr;  
    const Status status = queue.pop(item);  
    if (status != Status::SUCCESS) {  
        break;  
    }  
    dispatch(item);  
}
```

**Bounded Loop**

**Unbounded loops can run forever and hang a thread  
This behavior is bad in flight code**



# Avoid Recursion

- This is a form of a potentially unbounded loop
- It can also cause stack overflow

```
int32_t factorial(int32_t x) {  
    if (x <= 1) {  
        return 1;  
    }  
    else {  
        return x * factorial(x - 1);  
    }  
}
```

**Recursion**

```
int32_t factorial(int32_t x) {  
    int32_t result = 1;  
    for (int32_t i = 2; i <= x; i++) {  
        result *= i;  
    }  
    return result;  
}
```

**Bounded Loop**

**This code should also check for integer overflow**

# Outline

- Check numeric bounds
- Use bounded loops
- **Use assertions**
- Avoid dangerous code
- Run static analysis
- Follow a coding standard



# Use Assertions

- An **assertion** is a macro expression  $ASSERT(e)$ 
  - $e$  is a Boolean expression that is expected to evaluate to *true*
  - If  $e$  does evaluate to *true*, the macro does nothing
  - Otherwise an **assertion failure** occurs
- When an assertion failure occurs
  - In unit testing, we usually halt the program
  - In system test and flight
    - At JPL we generally restart, unless restart could cause mission failure
    - Other institutions turn off assertion responses (ignore the failure)

**Use of assertions is important for testing, debugging, and maintenance**



# When To Use Assertions

- Use assertions when failure “cannot happen” (e.g., it indicates a bug)
  - Pointer is null just before dereference
  - Array index is out of bounds just before array access
  - Computed value is outside expected range
- Otherwise check for and handle the error (e.g., bad user input)
- Common pattern:
  - Check for and reject invalid input
  - After passing the check, assert that input is valid before it is used
  - Assertion failure indicates a bug in the code that did the checking

**Never assert that ground or user input is valid**



# Outline

- Check numeric bounds
- Use bounded loops
- Use assertions
- **Avoid dangerous code**
- Run static analysis
- Follow a coding standard



# Avoid Dangerous Code

- C and C++ have many traps for the unwary
- It is easy to write and run programs that have **undefined** behavior
  - Relatively straightforward
    - Read of uninitialized memory
    - Object access through dangling pointer
  - More subtle
    - Left shift of negative signed integer (C, C++)
    - Array index starting at pointer to base-class object (C++)
- Some behavior is defined but obscure
  - Tricky operator precedence (use parentheses to disambiguate)
  - Complex C macros (avoid)
  - Forgetting to declare a base-class destructor *virtual* (C++) (don't do it)



**DANGER**

**C AND C++ HAZARDS**



# Avoid Complicated C macros

## defensive coding: testable code

example: make very limited use of the C preprocessor

Q1: will this code trigger  
a compilation error (and if so, where)?

```
#define A

int
main(void)
{
    #ifndef A
        printf("hello world\n");
    #else1
        goto *main();
    #endif
    return 0;
}
```

```
$ gcc -Wall -pedantic gobble.c
$ ./a.out
$
```

Q2: how many different ways can this  
code be compiled (i.e., how many  
ways would it need to be tested)?

```
#if (a>0)
...
    #ifdef X
    ...
        #ifndef Y
        ...
            #if b
            ...
                #endif
            ...
        #endif
    ...
    #endif
...
#endif
```

A: 16 different ways ( $2^4$ )

**Fewer tests may suffice,  
but correctness may be  
difficult to prove**





# Avoid Side Effects in Macro Arguments

```
#define ABS(X) (((X) < 0) ? -(X) : (X))
```

**Macro for computing  
absolute value**

```
int32_t a = get_a();  
a++;  
int32_t b = ABS(a);
```

**OK**

```
int32_t a = get_a();  
int32_t b = ABS(a++);
```

**Probably not what  
is intended**



# Avoid Side Effects in Index Expressions

```
#include <cinttypes>
#include <array>

int32_t strangeAddition(
    const std::array<int32_t, 10>& a,
    size_t i
){
    return a[i] + a[i++];
}
```

**Evaluation order is not  
guaranteed**



# Guard Against Out-of-Bounds Array Access

```
class History {  
    ...  
    public:  
        Item& getItemAt(size_t index) {  
            ASSERT(index < HISTORY_SIZE);  
            return m_items[index];  
        }  
    ...  
    private:  
        Item m_items[HISTORY_SIZE];  
};
```

- Use error checking or assertions to ensure that array accesses stay in bounds
- Otherwise out-of-bounds access can cause strange behavior that is hard to debug
- If performance is extremely critical, these checks can be disabled in system testing and flight



# Avoid Dangling Pointers

```
int32_t *stackReturn() {  
    int32_t x;  
    return &x;  
};
```

**x is out of scope on return  
from the function**

```
int32_t outOfScope() {  
    int32_t *p = nullptr;  
    {  
        int32_t x = 5;  
        p = &x;  
    }  
    return *p;  
}
```

**x is out of scope when the  
inner block ends**



# Avoid Dynamic Memory Allocation

- Allocate all memory you need when the program starts
- Don't use *malloc/new* or *free/delete* after initialization
- If you need dynamic behavior, use a memory pool
- This rule
  - Avoids heap fragmentation, which leads to nondeterministic performance
  - Avoids undefined behavior due to dangling pointers and memory leaks



# Outline

- Check numeric bounds
- Use bounded loops
- Use assertions
- Avoid dangerous code
- **Run static analysis**
- Follow a coding standard



# Run Static Analysis

- Compile with *--Wall --Werror*
- Run static analysis tools
  - Free tools, e.g., as GitHub actions
  - Commercial tools such as Coverity and CodeSonar
- Run the tools as often as feasible
  - Compiler: On every build
  - GitHub actions: On every pull request to main
  - Other tools: On every code review

# Outline

- Check numeric bounds
- Use bounded loops
- Use assertions
- Avoid dangerous code
- Run static analysis
- **Follow a coding standard**





# Follow a Coding Standard

- A **coding standard** is an agreed-upon set of rules for developing code
- Enforcement is usually by a combination of
  - Automated tools
  - Code review
- Examples
  - JPL C, C++
  - MISRA C, C++
  - AUTOSAR C++
  - Google C++ Style Guide

**It is important to agree upon and adhere to a coding standard**



# The Power of 10

## defensive coding: **coding standards**

follow a machine-checkable, risk-based, standard

1. *Restrict to simple control flow constructs*
2. *Do not use recursion and give all loops a fixed upper-bound*
3. *Do not use dynamic memory allocation after initialization*
4. *Limit functions to no more than ~60 lines of text*
5. *Target an average assertion density of 2% per module*
6. *Declare data objects at the smallest possible level of scope*
7. *Check the return value of non-void functions; check the validity of parameters*
8. *Limit the use of the preprocessor to file inclusion and simple macros*
9. *Limit the use of pointers. Use no more than 2 level of dereferencing*
10. *Compile with all warnings enabled, and use source code analyzers*

Gerard Holzmann



IEEE Computer 39(6) 95-97 (2006)

<http://spinroot.com/p10/>

# References

Power of 10 rules for C: <https://spinroot.com/p10>

The MISRA coding standards: <https://www.misra.org.uk>

The AUTOSAR coding standards: <https://www.autosar.org>

## Static analysis

- Overview of static analyzers: <https://spinroot.com/static>
- CodeSonar: <https://www.grammatech.com/codesonar-cc>
- Coverity: <https://www.synopsys.com/software-integrity.html>
- KlockWork: <https://www.perforce.com/products/klocwork>
- Semmle: <https://github.blog/2019-09-18-github-welcomes-semmle>
- Cobra: <https://spinroot.com/cobra> (free)
- Frama-C: <https://frama-c.com> (free)
- uno: <https://spinroot.com/uno> (free)

## Dynamic analysis

- Test driven development: [https://en.wikipedia.org/wiki/Test-driven\\_development](https://en.wikipedia.org/wiki/Test-driven_development)
- Fuzz testing: <https://www.fuzzingbook.org> (Andreas Zeller)
- Valgrind: <https://valgrind.org>
- Mutation testing: <https://github.com/mull-project/mull>